

```
# === 1) LOAD ===
import pandas as pd
import numpy as np

movies = pd.read_csv("movies.csv")
ratings = pd.read_csv("ratings.csv")

print(movies.shape) # e.g., (62423, 3)
print(ratings.shape) # e.g., (453712, 4)
display(movies.head())
display(ratings.head())

# --- Join ratings with movie text (title + genres) ---
mov = movies.copy()
mov["text"] = (mov["title"] .fillna("") + " | " + mov["genres"] .fillna("")).str.lower()
df_full = ratings.merge(mov[["movieId", "text", "title", "genres"]], on="movieId", how="inner")

# For per-user training we'll define like/dislike later (y) using >=4 as like, <=2.5 as dislike (drop neutrals)
```

(62423, 3)
(43149, 4)

	movieId	title	genres	
0	1	Toy Story (1995)	Adventure Animation Children Comedy Fantasy	
1	2	Jumanji (1995)	Adventure Children Fantasy	
2	3	Grumpier Old Men (1995)	Comedy Romance	
3	4	Waiting to Exhale (1995)	Comedy Drama Romance	
4	5	Father of the Bride Part II (1995)	Comedy	
	userId	movieId	rating	timestamp
0	1	296	5.0	1.147880e+09
1	1	306	3.5	1.147869e+09
2	1	307	5.0	1.147869e+09
3	1	665	5.0	1.147879e+09
4	1	899	3.5	1.147869e+09

```
# === 2) A HELPER to build a per-user dataset (like your spam PREPARE DATA) ===
from sklearn.model_selection import train_test_split
```

```
def make_user_dataset(user_id, neutral_cut=3.0):
    """Return X_train, X_test, y_train, y_test, pool_unseen for a given user.
    Likes: rating >= 4.0; Dislikes: rating <= 2.5; Drop neutrals (around 3.0).
    pool_unseen = all movies the user hasn't rated yet (to rank & recommend)."""
    user_df = df_full[df_full["userId"] == user_id].copy()
    if user_df.empty:
        raise ValueError(f"User {user_id} not found.")

    # label: 1 = like, 0 = dislike; drop neutrals near 3.0
    user_df = user_df.assign(
        y = np.where(user_df["rating"] >= 4.0, 1,
                    np.where(user_df["rating"] <= 2.5, 0, np.nan))
    ).dropna(subset=["y"])
    user_df["y"] = user_df["y"].astype(int)

    # Text features (title+genres string)
    X = user_df["text"].astype(str)
    y = user_df["y"].astype(int)

    # Keep a pool of unseen movies for recommendation
    seen_ids = set(user_df["movieId"].unique())
    pool_unseen = mov[~mov["movieId"].isin(seen_ids)].copy()

    # Guard: need at least some positives and negatives
    pos, neg = (y==1).sum(), (y==0).sum()
    if pos < 5 or neg < 5:
        print(f"⚠️ User {user_id} has limited labels (pos={pos}, neg={neg}). Results may be weak.")

    X_train, X_test, y_train, y_test = train_test_split(
        X, y, test_size=0.25, stratify=y, random_state=42
```

```

        )
    return X_train, X_test, y_train, y_test, pool_unseen

# === 3) PREPARE MODELS (very similar to your spam project) ===
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.linear_model import LogisticRegression, SGDClassifier
from sklearn.naive_bayes import MultinomialNB, ComplementNB
from sklearn.svm import LinearSVC
from sklearn.calibration import CalibratedClassifierCV

pipelines = {
    "LogisticRegression": Pipeline([
        ("tfidf", TfidfVectorizer(ngram_range=(1,2), min_df=2)),
        ("clf", LogisticRegression(max_iter=200, class_weight="balanced", random_state=42)),
    ]),
    "MultinomialNB": Pipeline([
        ("tfidf", TfidfVectorizer(ngram_range=(1,2), min_df=2)),
        ("clf", MultinomialNB()),
    ]),
    "ComplementNB": Pipeline([
        ("tfidf", TfidfVectorizer(ngram_range=(1,2), min_df=2)),
        ("clf", ComplementNB()),
    ]),
    # LinearSVC is strong; wrap for predict_proba
    "LinearSVC+Calibrated": Pipeline([
        ("tfidf", TfidfVectorizer(ngram_range=(1,2), min_df=2)),
        ("clf", CalibratedClassifierCV(LinearSVC(random_state=42), cv=3)),
    ]),
    "SGD (log loss)": Pipeline([
        ("tfidf", TfidfVectorizer(ngram_range=(1,2), min_df=2)),
        ("clf", SGDClassifier(loss="log_loss", class_weight="balanced", random_state=42)),
    ]),
}
}

# === 4) TRAIN & EVALUATE for a chosen user (like spam training loop) ===
from sklearn.metrics import accuracy_score, classification_report, confusion_matrix

user_id = 1 # <- change to any existing user Id
X_train, X_test, y_train, y_test, pool_unseen = make_user_dataset(user_id)

results = []
best_name, best_acc = None, -1.0
for name, pipe in pipelines.items():
    pipe.fit(X_train, y_train)
    pred = pipe.predict(X_test)
    acc = accuracy_score(y_test, pred)
    results.append((name, acc))
    if acc > best_acc:
        best_acc, best_name = acc, name
print(f"\n== {name} ==")
print("Accuracy:", round(acc, 3))
print(classification_report(y_test, pred, target_names=["dislike", "like"], digits=3))

print("\nSUMMARY:", sorted(results, key=lambda x: x[1], reverse=True))
best_pipe = pipelines[best_name]
print(f"\nBest model: {best_name} | Accuracy: {best_acc:.3f}")

cm = confusion_matrix(y_test, best_pipe.predict(X_test))
print("Confusion matrix [[TN, FP],[FN, TP]]:\n", cm)

import matplotlib.pyplot as plt
plt.figure(figsize=(4,4))
plt.imshow(cm, interpolation='nearest')
plt.title("Confusion Matrix")
plt.xticks([0,1], ["dislike", "like"])
plt.yticks([0,1], ["dislike", "like"])
for i in range(2):
    for j in range(2):
        plt.text(j, i, cm[i, j], ha="center", va="center")
plt.xlabel("Predicted")
plt.ylabel("Actual")
plt.tight_layout()
plt.show()

```



```
==== LogisticRegression ====
Accuracy: 0.615
      precision    recall   f1-score   support
dislike       0.000     0.000     0.000      3
like        0.727     0.800     0.762     10

accuracy      0.615
macro avg    0.364     0.400     0.381     13
weighted avg  0.559     0.615     0.586     13
```

```
==== MultinomialNB ====
Accuracy: 0.769
      precision    recall   f1-score   support
```

```
# === 5) RECOMMEND (score all unseen movies for this user and show Top-N) ===
def predict_proba_safe(pipe, texts):
    if hasattr(pipe, "predict_proba"):
        try:
            return pipe.predict_proba(texts)[:,1]
        except Exception:
            pass
    if hasattr(pipe, "decision_function"):
        from scipy.special import expit
        return expit(pipe.decision_function(texts))
    # fallback: hard predictions only -> map to {0,1} as "prob"
    preds = pipe.predict(texts)
    return preds.astype(float)

def recommend_for_user(pipe, user_id, pool_unseen, top_n=10):
    texts = pool_unseen["text"].astype(str)
    probs = predict_proba_safe(pipe, texts)
    pool_unseen = pool_unseen.assign(p_like=probs)
    top = pool_unseen.sort_values("p_like", ascending=False).head(top_n)
    return top[["title", "genres", "p_like"]]

top10 = recommend_for_user(best_pipe, user_id, pool_unseen, top_n=10)
display(top10)
```

		accuracy	precision	recall	f1-score	support	
		macro avg	0.385	0.500	0.435	13	
26000	The Helen Morgan Story (1957)	Drama Romance	0.945439				
25879	Sea Wife (1957)	Drama Romance	0.944567				
36293	Cinderella (1957)	Drama Romance	0.944567				
7232	Peyton Place (1957)	Drama Romance	0.944567				
31482	Bombers B-52 (1957)	Drama Romance	0.944567				
15155	Kisses (Kuchizuke) (1957)	Drama Romance	0.944567				
54751	Berlin, Schoenhauser Corner (1957)	Drama Romance	0.944567				
38248	Rot ist die Liebe (1957)	Drama Romance	0.944567				
11052	Notti bianche, Le (White Nights) (1957)	Drama Romance	0.944567				
11108	Raintree County (1957)	Drama Romance	0.944567				

Best model: MultinomialNB | Accuracy: 0.769
Confusion matrix: $\begin{bmatrix} TN & FP \\ FN & FN \end{bmatrix}$
Next steps: [[Generate code with top10](#)] [[New interactive sheet](#)]

```
# === 6) TINY GUI (like your spam GUI) ===
import ipywidgets as widgets
from IPython.display import display, clear_output

user_in = widgets.IntText(value=1, description="User ID:")
model_dd = widgets.Dropdown(options=list(pipelines.keys()), value=best_name, description="Model:")
train_btn = widgets.Button(description="(Re)Train", button_style="primary")
top_in = widgets.IntSlider(value=10, min=5, max=30, step=1, description="Top-N:")
pred_btn = widgets.Button(description="Recommend", button_style="")
out = widgets.Output(layout={'border': '1px solid #ddd'})

state = {"pipe": best_pipe, "name": best_name, "acc": best_acc,
         "X_train": X_train, "X_test": X_test, "y_train": y_train, "y_test": y_test,
         "pool": pool_unseen, "user": user_id}
```

```

def retrain_for_user(uid, name):
    X_tr, X_te, y_tr, y_te, pool = make_user_dataset(uid)
    pipe = pipelines[name]
    pipe.fit(X_tr, y_tr)
    acc = accuracy_score(y_te, pipe.predict(X_te))
    return pipe, acc, X_tr, X_te, y_tr, y_te, pool

def on_train_clicked(_):
    with out:
        clear_output(wait=True)
        uid = int(user_in.value)
        name = model_dd.value
        try:
            pipe, acc, X_tr, X_te, y_tr, y_te, pool = retrain_for_user(uid, name)
            state.update({"pipe": pipe, "name": name, "acc": acc,
                          "X_train": X_tr, "X_test": X_te, "y_train": y_tr, "y_test": y_te,
                          "pool": pool, "user": uid})
            print(f"✓ Trained {name} for User {uid} | Test Acc: {acc:.3f} | Train size: {X_tr.shape[0]}")
        except Exception as e:
            print("Error:", e)

def on_recommend_clicked(_):
    with out:
        if state["pipe"] is None:
            print("Please train a model first.")
            return
        topn = int(topn_in.value)
        recs = recommend_for_user(state["pipe"], state["user"], state["pool"], top_n=topn)
        print(f"Top-{topn} recommendations for User {state['user']} (Model: {state['name']}, Acc {state['acc']:.3f})")
        display(recs)

train_btn.on_click(on_train_clicked)
pred_btn.on_click(on_recommend_clicked)

ui = widgets.VBox([
    widgets.HBox([user_in, model_dd, train_btn]),
    widgets.HBox([topn_in, pred_btn]),
    out
])
display(ui)

with out:
    print(f"Ready. Current user={state['user']} | Model={state['name']} (Acc {state['acc']:.3f}). Click Recommend.")

```

User ID: Model: (Re)Train

Top-N: Recommend

Trained MultinomialNB for User 1 | Test Acc: 0.769 | Train size: 36